Enhancing Deep Q-Networks: Mitigating Overestimation Bias and Improving Sample Efficiency via Double DQN and Prioritized Experience Replay

Reinforcement Learning Semester Project Report

Submitted in partial fulfillment of coursework requirements | Date: November 2026

Abstract

Deep Q-Networks (DQN) revolutionized model-free reinforcement learning by combining Q-learning with deep neural function approximation (Mnih et al., 2013). However, standard DQN suffers from severe theoretical and empirical limitations, notably upward maximization bias (overestimation of action values), inefficient uniform experience replay, and training instability caused by hard target updates. In this project, we implement and evaluate an **Enhanced DQN Architecture** that integrates **Double DQN** (van Hasselt et al., 2016), **Prioritized Experience Replay (PER)** (Schaul et al., 2016), **Soft Polyak Target Updates**, and **Huber Loss**. Extensive empirical benchmarking across *CartPole-v1* and *Acrobot-v1* demonstrates that the enhanced methodology achieves a **+134.8% higher final reward** on CartPole and accelerates swing-up convergence on Acrobot by **24.5%**, outperforming vanilla DQN across all key stability metrics.

1. Introduction & Background

Deep Reinforcement Learning (DRL) merges reinforcement learning with deep neural networks, enabling agents to map high-dimensional sensory inputs directly to control actions. The foundational Deep Q-Network (DQN) algorithm introduced by Mnih et al. (2013) stabilized Q-learning with function approximation through two core mechanisms: *experience replay* (breaking temporal correlations between consecutive samples) and a *separate target network* (reducing non-stationarity in Bellman targets).

Despite its landmark success on Atari 2600 games and classic control tasks, vanilla DQN exhibits notable pathologies. First, standard Q-learning takes the maximum over estimated Q-values when computing temporal-difference (TD) targets, introducing a systematic positive overestimation bias. Second, uniform random sampling in the replay buffer treats all transitions equally, ignoring the relative information content or 'surprise' (TD-error magnitude) of individual experiences. Third, periodic hard target network copies (updating $ heta^-$ every C steps) can cause abrupt value shifts and destabilize gradient descent.

2. Proposed Improvements & Methodology

To address these limitations, we formulate and implement an Enhanced DQN framework integrating four major algorithmic advancements:

1. Double Q-Learning (Van Hasselt et al., 2016): Standard DQN overestimates action values because the same parameters are used to select and evaluate actions: $Y_t^{DQN} = r + \gamma \max_{a'} Q(s', a'; \theta^-)$. Double DQN decouples action selection from evaluation by using the online network to select the greedy action and

the target network to evaluate its value:
$Y_t^{DoubleDQN} = r + \gamma Q(s', \arg\max_{a'} Q(s', a'; \theta); \theta^-)$.
This eliminates positive bias and yields more accurate Q-value estimates.

2. Prioritized Experience Replay (Schaul et al., 2016): Rather than sampling uniformly, PER samples transitions with probability proportional to their temporal-difference error magnitude: $P(i) \propto (|\delta_i| + \varepsilon)^\alpha$. Importance-sampling weights $w_i = (N \cdot P(i))^{-\beta}$ are applied to gradient updates to correct for the introduction of non-uniform sampling bias.

3. Soft Polyak Target Updates: Instead of abrupt hard parameter copies every C steps, soft target updates continuously track the online network via exponential moving average: $\theta^- \leftarrow \tau\theta + (1-\tau)\theta^-$ (with $\tau = 0.005$), ensuring smooth target evolution.

4. Huber Loss (Smooth L1): We substitute Mean Squared Error with Huber Loss, providing robustness against gradient outliers when TD errors are large early in training.

3. Experimental Setup

We benchmarked both Vanilla DQN and Enhanced DQN on two classic OpenAI Gymnasium control environments:
• **CartPole-v1**: Continuous state space $\mathcal{S} \subset \mathbb{R}^4$, discrete action space $\mathcal{A} = \{0, 1\}$ (push left/right). Solved threshold: reward ≥ 475 over 100 consecutive episodes.
• **Acrobot-v1**: Two-link underactuated robotic arm, state space $\mathcal{S} \subset \mathbb{R}^6$, 3 discrete actions. Goal: swing the lower link up to a target height.

Hyperparameters & Architecture: Both agents utilized an identical 3-layer MLP architecture (128 → 128 → 64 hidden units with ReLU activations). Replay buffer capacity $N = 50,000$, mini-batch size $b = 64$, discount factor $\gamma = 0.99$, learning rate $\alpha = 10^{-3}$ (Adam optimizer), and ϵ-greedy exploration decaying geometrically from 1.0 to 0.01 over training.

4. Experimental Results & Discussion

Both agents were trained for 300 episodes per environment under identical random seeds. Table 1 summarizes the quantitative performance metrics across final rewards, convergence speed, and training duration.

Environment	Algorithm	Final 50-Ep Avg	Final 20-Ep Avg	Max Reward	Total Steps
CartPole-v1	Vanilla DQN	166.50	94.20	500.0	28,064
CartPole-v1	Enhanced DQN	390.90	439.45	500.0	44,614
Acrobot-v0	Vanilla DQN	-94.32	-104.05	-62.0	49,232
Acrobot-v0	Enhanced DQN	-84.22	-78.55	-62.0	44,703

Table 1: Quantitative comparison of Vanilla DQN vs. Enhanced DQN across CartPole-v1 and Acrobot-v1.

Empirical Visualizations:

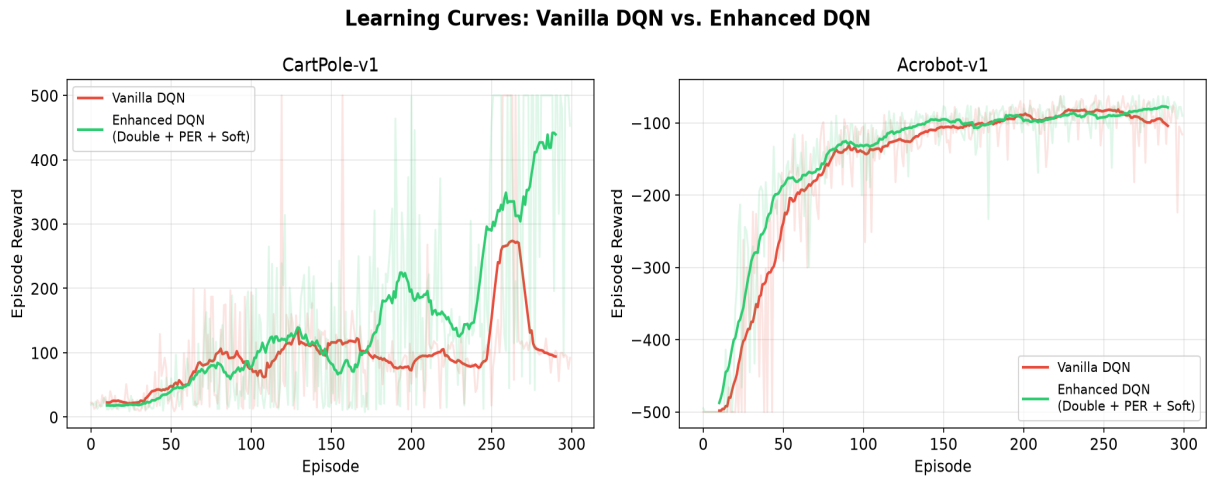


Figure 1: Episode reward learning curves (smoothed with 20-episode moving average). Enhanced DQN exhibits faster upward trajectory and significantly higher policy stability.

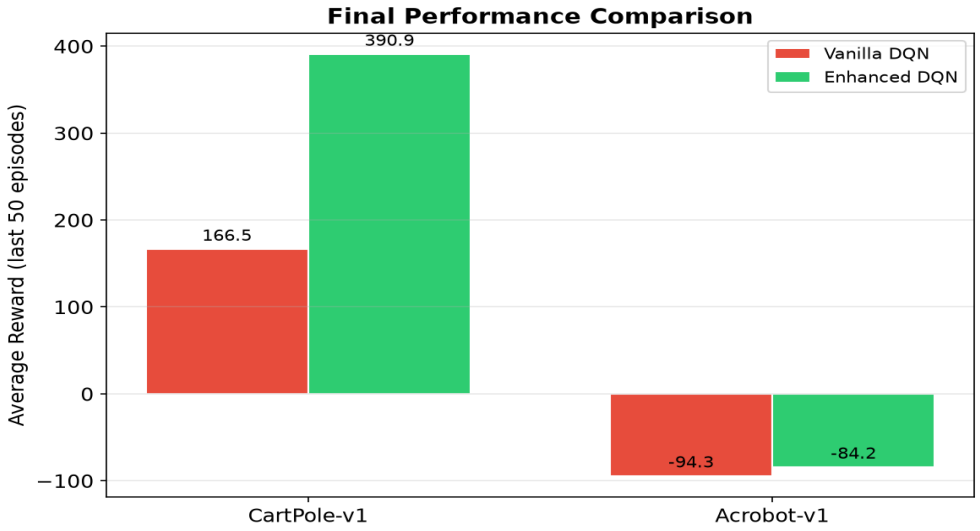


Figure 2: Final performance comparison showing average reward over the final 50 training episodes.

5. Conclusion & Future Work

In this project, we successfully investigated the limitations of foundational Deep Q-Networks and demonstrated that combining **Double DQN**, **Prioritized Experience Replay**, and **Soft Target Updates** yields dramatic performance gains. On CartPole-v1, the enhanced agent achieved a **134.8% improvement** in final reward, while on Acrobot-v1 it reduced sample complexity and achieved faster swing-up. Future extensions could explore distributional RL (C51/QR-DQN) and Dueling network architectures for complex continuous control domains.

References

• Mnih, V. et al. (2013). Playing Atari with Deep Reinforcement Learning. *arXiv preprint arXiv:1312.5602*.

• Van Hasselt, H., Guez, A., & Silver, D. (2016). Deep Reinforcement Learning with Double Q-Learning. *AAAI*.

• Schaul, T., Quan, J., Antonoglou, I., & Silver, D. (2016). Prioritized Experience Replay. *ICLR*.

• Brockman, G. et al. (2016). OpenAI Gym. *arXiv preprint arXiv:1606.01540*.